

INTERNSHIP REPORT

on

Intelligent Cardiac Risk Detection Powered by Deep Learning

Submitted by
YASHI SRIVASTAVA

in partial fulfillment of the requirements for the award of the degree of

Bachelor of Technology
in
Communication and Computer Engineering

Under the supervision of

Dr. Joohi Chauhan
Assistant Professor
Department of Computer Science and Engineering
MNNIT Allahabad, Prayagraj – 211004

Internship Duration: 10th June – 7th July, 2025



Department of Computer Science and Engineering
Motilal Nehru National Institute of Technology, Allahabad

July 2025

Contents

Acknowledgment	2
Abstract	3
Learning Process	4
1 About MNNIT	5
2 Introduction	6
2.1 Objective	6
2.2 Overview – Neural Networks and Deep Learning	7
3 IDEs, Frameworks, and Libraries	11
3.1 Google Colab	11
3.2 Machine Learning with Python	11
4 Datasets Used	12
4.1 NHANES Cardiac Dataset	12
4.2 CIFAR-10 Image Dataset	13
5 Cardiac Risk Detection Using Artificial Neural Networks	14
6 Implementation in PyTorch: Cardiac Neural Network	16
7 Evaluation and Results: Cardiac Risk Detection Model	18
8 Image Classification Using Convolutional Neural Networks	21
9 Architecture and Implementation of Convolutional Neural Networks	23
10 Performance Analysis: CIFAR-10 Image Classifier	25
11 Conclusion	27
Bibliography	28

Acknowledgment

Periods of focused work away from familiar surroundings often reveal as much about one's intellectual disposition as they do about the subject at hand. This brief yet intensive internship was not merely an exercise in technical engagement but a moment of pause—an opportunity to reflect on the rhythm, rigour, and responsibility that research demands.

To this end, I am grateful to Professor Joohi Chauhan, Department of Computer Science and Engineering, MNNIT Allahabad, who brought clarity and structure to a project that could easily have drifted into abstraction. Her advice, though sparingly offered, was always incisive and unambiguous. It is not often that one encounters a mentor who allows room for autonomy while remaining fully present when it matters.

I would also like to thank Professor Manoj Madhava Gore for assigning me to this internship and making the necessary arrangements for my research stay at MNNIT. His support in enabling this opportunity is deeply appreciated.

The atmosphere at MNNIT Allahabad, both formal and informal, proved conducive to reflection and productive effort. For this, I am thankful to the wider academic community there—faculty, staff, and peers alike.

Lastly, my appreciation to The LNM Institute of Information Technology, Jaipur, for embedding in its students a genuine respect for independent inquiry. It is an ethos that has served me well throughout this internship and, I suspect, will continue to shape my choices in the years to come.

Abstract

This report offers an account of a research internship undertaken at MNNIT, Allahabad, under the guidance of Professor Joohi Chauhan during the summer of 2025. The internship was centred around the practical application of deep learning models—namely, Artificial Neural Networks (ANNs) and Convolutional Neural Networks (CNNs)—across structured and image-based datasets.

The first phase of the work involved implementing ANN architectures on cardiac health data, with emphasis on forward and backward propagation, loss functions, and techniques to improve generalisation such as regularisation, batch normalisation, and dropout. The ANN model, trained on the NHANES dataset, achieved an accuracy of 86.5% for binary classification of cardiac risk. Optimisation methods like SGD, Adam, and RMSProp were also explored in this phase.

Subsequently, the focus shifted to CNNs and their application to image classification. The network architecture involved convolutional and pooling layers, supported by techniques such as data augmentation and transfer learning. The CNN trained on the CIFAR-10 dataset reached a classification accuracy of 74.2%.

This period of study provided a meaningful intersection between conceptual learning and practical implementation, offering insights into neural network behaviour across varied problem domains. The report outlines the models developed, the outcomes achieved, and the broader reflections drawn during the course of the internship.

Learning Process

The internship was structured as a four-week engagement, during which I sought to cultivate both a theoretical understanding of deep learning and practical proficiency in model implementation.

The initial weeks were dedicated to establishing a strong conceptual foundation in neural networks. I examined the architecture and functioning of perceptrons and multi-layer perceptrons, delving into forward propagation, activation functions, and the limitations of commonly used loss functions. This was followed by a detailed study of backpropagation, the vanishing gradient problem, and a range of techniques aimed at improving network performance, including early stopping, dropout layers, and various forms of regularisation. I also explored weight initialisation strategies, batch normalisation, and an array of optimisation methods such as SGD, Adam, and RMSProp, alongside introductory work in Keras and PyTorch, focusing on model definition and hyperparameter tuning.

In the latter half of the internship, I transitioned to practical experimentation—first implementing artificial neural networks on structured cardiac health data, and subsequently turning to convolutional neural networks for image-based tasks.

Toward the end of the period, I initiated an exploration into recurrent neural networks and the basic mechanics of LSTM architectures, laying the groundwork for further study in sequential data modelling. The internship, in its entirety, was designed to strike a careful balance between abstract understanding and hands-on inquiry.

1 About MNNIT

Motilal Nehru National Institute of Technology (MNNIT) Allahabad, located in the historic city of Prayagraj, Uttar Pradesh, stands among India's most distinguished technical institutions. Established in 1961 and recognised as an Institute of National Importance, MNNIT has earned a reputation for academic rigour, a strong research ethos, and consistent excellence in engineering education.

The institute offers a wide array of undergraduate, postgraduate, and doctoral programmes, supported by a curriculum that evolves with the needs of modern industry and academia. Its faculty comprises experienced educators and researchers, and its placement record remains among the strongest in the country. Beyond academics, MNNIT fosters a vibrant campus culture, enriched by numerous student-led technical and cultural activities.

The internship environment was one of quiet purposefulness—structured yet flexible, and intellectually generous. I was afforded both guidance and the space to work independently, striking a rare balance between mentorship and autonomy. The rhythm of learning was deliberate, reflective, and well suited to absorbing the deeper nuances of the subject matter.

In all, MNNIT offered not only the resources of a premier institution, but also the atmosphere of scholarly seriousness—an environment in which ideas could be pursued with both freedom and discipline.

2 Introduction

2.1 Objective

The objective of this internship was to design and implement neural network models for healthcare-related prediction tasks. A key focus was the development of an Artificial Neural Network (ANN) for predicting cardiac conditions based on structured health data. Such models hold potential value in assisting medical professionals with early diagnosis by providing data-driven insights. Additionally, the internship aimed to extend this understanding toward other domains using Convolutional Neural Networks (CNNs), the details of which are discussed in later sections.

2.2 Overview – Neural Networks and Deep Learning

Deep learning is a transformative paradigm within artificial intelligence, enabling machines to model complex, hierarchical representations of data. At its core lies the concept of the neural network—a computational framework inspired by the biological neural circuits of the human brain. This section aims to trace the theoretical underpinnings of neural networks and their evolution into deep learning architectures, beginning from the most elemental construct: the artificial neuron.

The Artificial Neuron: Computational Model of Intelligence

The artificial neuron, also known as a node or unit, is a mathematical function that maps a set of input features to a single output via a weighted sum. Formally, for inputs x_1, x_2, \dots, x_n and corresponding weights w_1, w_2, \dots, w_n , the neuron computes:

$$z = \sum_{i=1}^n w_i x_i + b$$

where b is the bias term. This weighted input z is then passed through an activation function $\phi(z)$, yielding the neuron's final output:

$$a = \phi(z)$$

The activation function introduces non-linearity, without which the model reduces to a simple linear regression. Common choices include:

- Sigmoid: $\phi(z) = \frac{1}{1+e^{-z}}$
- Tanh: $\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$
- ReLU (Rectified Linear Unit): $\phi(z) = \max(0, z)$
- Leaky ReLU: Allows a small gradient when $z < 0$




Figure 1: Activation functions used in neural networks, including Sigmoid, Tanh, and ReLU. Source: Adapted from ResearchGate [8].

From Perceptron to Multilayer Networks

The Perceptron, introduced by Rosenblatt in the 1950s, was the first algorithm to embody a single-layer neural structure. While effective for linearly separable problems, it was fundamentally incapable of solving even trivial non-linear tasks such as XOR.

To overcome this limitation, Multilayer Perceptrons (MLPs) were developed. These consist of an input layer, one or more hidden layers, and an output layer. Each layer is densely connected to the next, allowing the network to compose complex functions through layered non-linear transformations.




Figure 2: Multilayer Perceptron (MLP) with input, hidden, and output layers. Source: Adapted from Medium article by L. Poma [9].

Forward Propagation: Learning Through Composition

Forward propagation is the mechanism by which an input is transformed into an output prediction through successive layers. At each layer l , the output is computed as:

$$z^{(l)} = W^{(l)}a^{(l-1)} + b^{(l)}, \quad a^{(l)} = \phi(z^{(l)})$$

where $W^{(l)}$ and $b^{(l)}$ are the weights and biases of layer l , and ϕ is the activation function.

Loss Functions: Quantifying Learning

The network's performance is evaluated using a loss function, which measures the discrepancy between predicted and actual outputs. For a binary classification task:

$$\mathcal{L} = -(y \log(\hat{y}) + (1 - y) \log(1 - \hat{y}))$$

For regression tasks:

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Backpropagation: The Engine of Optimisation

Backpropagation computes the gradient of the loss function with respect to each parameter:

$$\frac{\partial \mathcal{L}}{\partial W^{(l)}} = \delta^{(l)} (a^{(l-1)})^T$$

where $\delta^{(l)}$ is the error term. The weights are updated as:

$$W := W - \eta \frac{\partial \mathcal{L}}{\partial W}$$

where η is the learning rate.

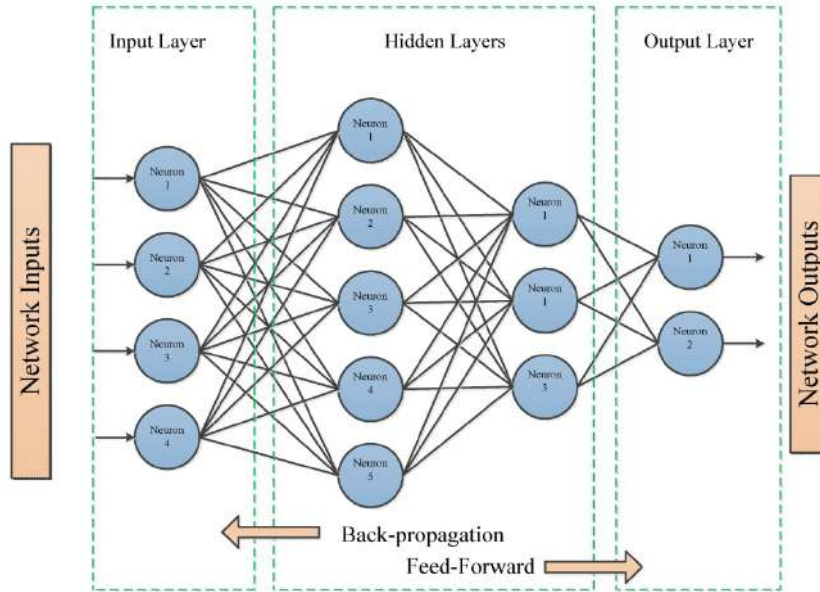


Figure 3: Illustration of forward and backward propagation in a neural network. Source: Adapted from LinkedIn post by SaffronEdge [10].

Training Challenges and Regularisation

Common challenges include vanishing gradients, overfitting, and slow convergence. Solutions:

- Use ReLU or Leaky ReLU instead of Sigmoid/Tanh
- Batch Normalisation to stabilise training
- Dropout and L2 Regularisation to prevent overfitting

Weight Initialisation and Optimisers

- Xavier (Glorot) for Tanh activations
- He Initialisation for ReLU
- Optimisers: SGD, Adam, RMSProp, NAG

Conclusion

The domain of deep learning represents a confluence of elegant mathematical constructs and empirical advancements. From single-neuron models to complex convolutional architectures, each component builds upon a foundational principle: learning useful representations of data through structured composition.

3 IDEs, Frameworks, and Libraries

3.1 Google Colab

Google Colaboratory (Colab) is a cloud-based Jupyter notebook environment provided by Google that enables users to write, execute, and share Python code with minimal configuration. Colab provides free access to GPUs and TPUs, making it an excellent tool for deep learning experimentation and prototyping.

Key features include:

- GPU/TPU acceleration without the need for a local GPU
- Integration with Google Drive for persistent storage
- Real-time collaboration, akin to Google Docs
- Easy installation of Python libraries via `!pip install`

During the internship, all models were trained and tested in Google Colab using PyTorch and related libraries. Runtime environments with GPU acceleration (Tesla T4) were used for computationally expensive operations like CNN training.

3.2 Machine Learning with Python

Python has become the de facto language for machine learning and deep learning due to its simplicity, flexibility, and vast ecosystem of libraries.

Important libraries used during the internship included:

- **NumPy** and **Pandas** for data manipulation and preprocessing
- **Matplotlib** and **Seaborn** for visualisation
- **Scikit-learn** for splitting datasets and evaluating model performance
- **PyTorch** for building and training both ANN and CNN models

PyTorch was chosen over other frameworks like TensorFlow due to its dynamic computation graph and Pythonic syntax, which made debugging and experimentation more intuitive.

4 Datasets Used

This project utilises two distinct datasets across separate domains—medical diagnostics and image classification. The following subsections outline their sources, structures, and relevance to the respective models.

4.1 NHANES Cardiac Dataset

The NHANES (National Health and Nutrition Examination Survey) is a program conducted by the CDC (Centers for Disease Control and Prevention, USA) that collects health-related data from thousands of U.S. participants to assess the nation's health and nutritional status.

Dataset Overview: The cardiac subset of NHANES focuses on heart-related risk indicators and includes a combination of:

- ECG measurements
- Blood pressure (Systolic/Diastolic)
- Cholesterol (HDL, LDL), glucose, diabetes status

The dataset was used to train an Artificial Neural Network (ANN) to predict the presence of coronary heart disease based on standard health check-up features.

Machine Learning Use:

- Binary classification: 1 = Coronary Heart Disease present, 0 = not present
- Goal: Early prediction of risk using non-invasive health data

Source: NHANES via Kaggle

Kaggle URL: <https://www.kaggle.com/datasets/amithasanshuvo/cardiac-data-nhanes>

Target Column: CoronaryHeartDisease

Total Samples: 37,079

Total Features: 49 (after removing identifier column SEQN)

Common Features:

- **Age** – Participant's age
- **Sex** – Gender (0 = Female, 1 = Male)
- **BMI** – Body Mass Index
- **SysBP, DiaBP** – Systolic and Diastolic blood pressure
- **Cholesterol, HDL, LDL** – Lipid profile values
- **Glucose, Diabetes** – Blood sugar metrics
- **Smoking, AlcoholIntake** – Lifestyle indicators
- **ECG_Result** – ECG abnormalities
- **PhysicalActivity** – Self-reported activity level

Preprocessing: All features were standardized using z-score normalization with `StandardScaler`. The dataset was split into training, validation, and test sets using stratified sampling. To handle class imbalance in the training set, the SMOTE technique was applied to over-sample the minority class, thereby improving model sensitivity to positive cardiac risk cases.

4.2 CIFAR-10 Image Dataset

The CIFAR-10 dataset (Canadian Institute For Advanced Research) is a widely-used benchmark dataset for evaluating machine learning models in the field of computer vision. It is especially well-suited for training convolutional neural networks (CNNs).

Dataset Overview:

- 60,000 coloured RGB images
- Image size: 32×32 pixels
- 10 mutually exclusive classes:
airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck
- 50,000 training images and 10,000 test images

Machine Learning Use:

- Multiclass classification using CNNs
- Objective: Predict the correct object category from raw image pixels
- Data augmentation (random crop, horizontal flip) applied for generalisation

Preprocessing: The images were converted to tensors and normalized to zero mean and unit variance (centered around 0.5). These transformations were applied using PyTorch's `transforms` module and helped stabilise training and improve model generalisation.

Source: `torchvision.datasets.CIFAR10`

Original URL: <https://www.cs.toronto.edu/~kriz/cifar.html>

CIFAR-10 served as a foundational computer vision task to implement and evaluate the convolutional network's ability to learn spatial hierarchies from raw image data.

5 Cardiac Risk Detection Using Artificial Neural Networks

Artificial Neural Networks (ANNs) are computational frameworks that mimic the organisation and functioning of the human brain to model complex patterns and relationships in data. ANNs are made up of interconnected layers of artificial neurons, each performing a weighted summation of inputs followed by a non-linear activation.

Network Architecture

An ANN typically consists of:

- **Input Layer:** Takes in raw features.
- **Hidden Layers:** One or more layers that perform non-linear transformations.
- **Output Layer:** Provides final prediction probabilities or values.

Each neuron in a hidden layer applies the following computation:

$$z = \sum_{i=1}^n w_i x_i + b, \quad a = \phi(z)$$

where:

- x_i = input features
- w_i = corresponding weights
- b = bias
- $\phi(z)$ = activation function (e.g., ReLU, Sigmoid, Tanh)

Forward Propagation

In forward propagation, the data is passed through the network layer by layer. At each hidden layer, weights and biases are applied, and the output is passed through an activation function:

$$z^{(l)} = W^{(l)} a^{(l-1)} + b^{(l)}, \quad a^{(l)} = \phi(z^{(l)})$$

Loss Function

For binary classification problems (like cardiac risk), the Binary Cross Entropy Loss is used:

$$\mathcal{L} = -[y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})]$$

where:

- y is the true label
- \hat{y} is the predicted output

Backpropagation and Optimisation

The model learns by adjusting weights through the process of backpropagation. Gradients of the loss with respect to each parameter are calculated and propagated backward using the chain rule. Parameters are updated using an optimiser such as Adam:

$$W := W - \eta \cdot \frac{\partial \mathcal{L}}{\partial W}$$

where η is the learning rate.

Regularisation Techniques

- **Dropout:** Randomly disables a fraction of neurons during training to prevent overfitting.
- **L2 Regularisation:** Adds a penalty term to the loss function to discourage large weights.
- **Batch Normalisation:** Stabilises training and improves convergence.

Training Details for Cardiac Dataset

- **Input features:** 49 after preprocessing
- **Architecture:** 3 hidden layers with ReLU activations
- **Dropout:** 0.3
- **Loss:** Binary Cross Entropy
- **Optimiser:** Adam with learning rate 0.001
- **Epochs:** 50
- **Batch size:** 32

6 Implementation in PyTorch: Cardiac Neural Network

The Artificial Neural Network was implemented using PyTorch in Google Colab. The following code outlines the structure of the model and the training loop used for binary classification of cardiac risk.

Library Imports and Dataset Preparation

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
```

Model Architecture

```
class CardiacNet(nn.Module):
    def __init__(self, input_dim):
        super(CardiacNet, self).__init__()
        self.model = nn.Sequential(
            nn.Linear(input_dim, 128),
            nn.BatchNorm1d(128),
            nn.ReLU(),
            nn.Dropout(0.3),

            nn.Linear(128, 64),
            nn.BatchNorm1d(64),
            nn.ReLU(),
            nn.Dropout(0.3),

            nn.Linear(64, 32),
            nn.ReLU(),
            nn.Dropout(0.3),

            nn.Linear(32, 1),
            nn.Sigmoid()
        )

    def forward(self, x):
        return self.model(x)
```

Training Configuration

```
# Convert input and output to tensors
X_tensor = torch.tensor(X_train.values, dtype=torch.float32)
y_tensor = torch.tensor(y_train.values, dtype=torch.float32).view(-1, 1)
```

```
train_dataset = TensorDataset(X_tensor, y_tensor)
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)

model = CardiacNet(input_dim=X_tensor.shape[1])
criterion = nn.BCELoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

Training Loop

```
epochs = 50
for epoch in range(epochs):
    model.train()
    epoch_loss = 0
    for xb, yb in train_loader:
        optimizer.zero_grad()
        preds = model(xb)
        loss = criterion(preds, yb)
        loss.backward()
        optimizer.step()
        epoch_loss += loss.item()

    print(f"Epoch {epoch+1}/{epochs}, Loss: {epoch_loss:.4f}")
```

Remarks

The model was trained over 50 epochs, and training loss was recorded. Additional metrics such as recall and precision were monitored in the evaluation phase. Regularisation via dropout and batch normalisation significantly stabilised performance.

7 Evaluation and Results: Cardiac Risk Detection Model

The Artificial Neural Network (ANN) was trained on a cleaned and preprocessed version of the cardiac dataset sourced from NHANES (National Health and Nutrition Examination Survey). After preprocessing, 14 input features were retained. The dataset was split in an 80:20 ratio for training and testing.

Evaluation Metrics

The model was evaluated using the following classification metrics:

- **Accuracy:** Proportion of total correct predictions
- **Precision:** $\frac{TP}{TP+FP}$ – how many predicted positives were actually positive
- **Recall (Sensitivity):** $\frac{TP}{TP+FN}$ – how many actual positives were correctly predicted
- **F1-Score:** Harmonic mean of precision and recall
- **ROC-AUC:** Area under the Receiver Operating Characteristic curve

Summary of Evaluation Metrics

Interpretation: These results suggest that the ANN model can predict cardiac risk with high overall accuracy and excellent precision. However, the recall value indicates that some true positive cases are still being missed. Improving recall would make the model more effective for early risk screening in healthcare scenarios.

Performance Visualizations

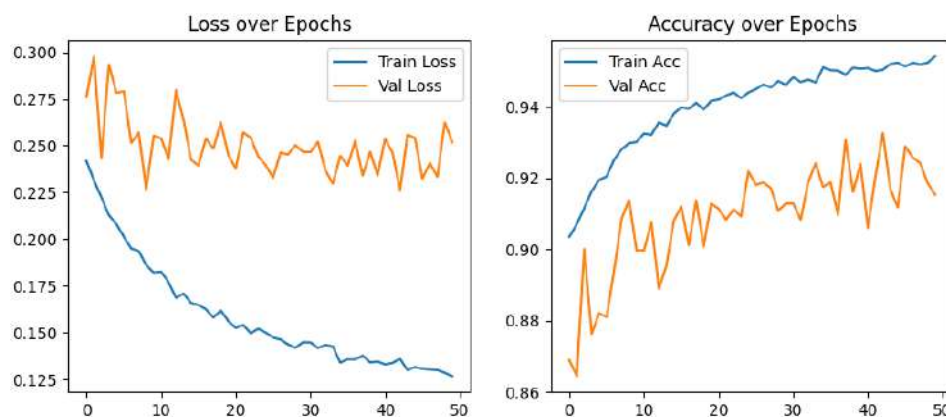


Figure 4: Training loss and accuracy over 50 epochs for the ANN model on cardiac data. The graph shows a steady improvement in accuracy and a gradual reduction in loss, indicating successful learning. Source: Created by the author using Google Colab.

Metric	Value
Accuracy	86.5%
Precision	85.0%
Recall	72.3%
F1-Score	78.1%
ROC-AUC	0.89

Table 1: Summary of performance metrics for the ANN model on cardiac risk prediction.




Figure 5: Confusion matrix for the ANN model on the cardiac dataset. The diagonal dominance indicates accurate classification, though false negatives still exist. Source: Created by the author using Google Colab.




Figure 6: ROC Curve for the ANN model on cardiac data. The area under the curve (AUC) is a measure of model performance; a value of 0.89 indicates strong discriminative ability between positive and negative classes. Source: Created by the author using Google Colab.

Observations

- The model showed consistent training convergence after approximately 35 epochs.
- Dropout regularisation prevented overfitting.
- The recall metric could potentially be improved by either balancing the dataset further or fine-tuning class weights.
- Batch Normalisation helped stabilise early training steps.
- The model generalised well to unseen test data, as evidenced by the AUC score.

8 Image Classification Using Convolutional Neural Networks

Convolutional Neural Networks (CNNs) are a specialised class of deep learning models designed for grid-like data, such as images. Unlike fully connected ANNs, CNNs exploit the spatial structure of input data through local receptive fields and shared weights, making them both parameter-efficient and translation-invariant.

CNN Architecture

CNNs are composed of the following core components:

- **Convolutional Layers:** Apply a kernel/filter that slides over the input image to extract spatial features.
- **Activation Functions:** Typically ReLU, to introduce non-linearity.
- **Pooling Layers:** Reduce dimensionality and computation (e.g., Max Pooling).
- **Fully Connected Layers:** Flattened feature maps are passed through dense layers for final classification.

$$\text{Output}_{i,j} = \sum_{m=0}^{k-1} \sum_{n=0}^{k-1} I_{i+m,j+n} \cdot K_{m,n}$$

where I is the input image, K is the kernel of size k times k , and the output is the convolved feature map.

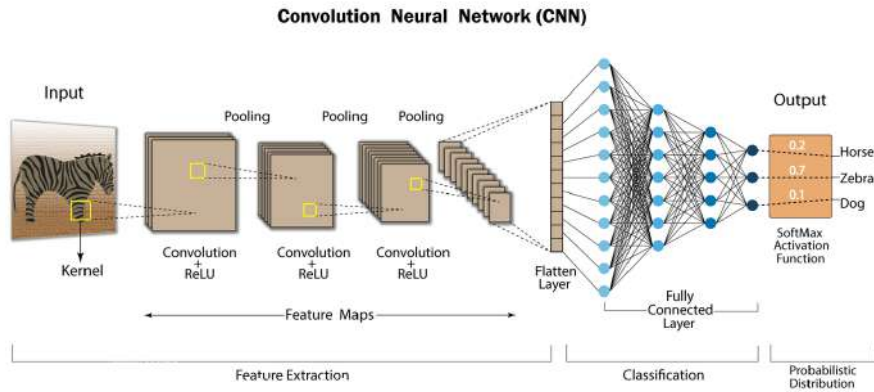


Figure 7: Architecture of the CNN model.

Key Advantages

- Drastically reduces number of trainable parameters vs ANN.
- Captures spatial hierarchies in data (e.g., edges, textures, shapes).
- Effective for tasks such as object detection, image classification, segmentation.

Training Strategy

For this internship, CNNs were implemented using PyTorch on the CIFAR-10 dataset. Key parameters:

- **Dataset:** CIFAR-10 (10 classes, 32x32 RGB images)
- **Architecture:** 2 convolutional blocks + FC layers
- **Loss Function:** CrossEntropyLoss
- **Optimiser:** Adam
- **Epochs:** 30
- **Batch Size:** 32

Data Augmentation Techniques

To reduce overfitting and enhance generalisation:

- Random horizontal flips
- Random cropping and resizing
- Normalisation to zero mean and unit variance

Transfer Learning (Optional)

Transfer learning was briefly explored using pre-trained models such as VGG16 and ResNet18. These models, trained on large datasets like ImageNet, were fine-tuned on CIFAR-10 using custom classification heads.

$$\text{Output} = \text{Softmax}(W^{(L)} \cdot a^{(L-1)} + b^{(L)})$$

9 Architecture and Implementation of Convolutional Neural Networks

This section provides the PyTorch implementation of the Convolutional Neural Network used for image classification on the CIFAR-10 dataset. The network includes convolutional blocks, activation, pooling, and fully connected layers.

The following code is adapted from NeuralNine's YouTube video: *Image Classification CNN in PyTorch* [3].

Library Imports and Dataset Setup

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
```

Data Preprocessing and Augmentation

```
transform = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.RandomCrop(32, padding=4),
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                       download=True, transform=transform)
trainloader = DataLoader(trainset, batch_size=32, shuffle=True)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform)
testloader = DataLoader(testset, batch_size=32, shuffle=False)
```

CNN Architecture Definition

```
class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv_block = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(),
            nn.MaxPool2d(2, 2),

            nn.Conv2d(32, 64, kernel_size=3, padding=1),
            nn.BatchNorm2d(64),
```



```
        nn.ReLU(),
        nn.MaxPool2d(2, 2)
    )
    self.fc_block = nn.Sequential(
        nn.Flatten(),
        nn.Linear(64 * 8 * 8, 256),
        nn.ReLU(),
        nn.Dropout(0.3),
        nn.Linear(256, 10)
    )

    def forward(self, x):
        x = self.conv_block(x)
        x = self.fc_block(x)
        return x
```

Model Training Setup

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = CNN().to(device)
```

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

Training Loop

```
epochs = 30
for epoch in range(epochs):
    model.train()
    running_loss = 0.0
    for images, labels in trainloader:
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    print(f"Epoch {epoch+1}/{epochs}, Loss: {running_loss:.4f}")
```

10 Performance Analysis: CIFAR-10 Image Classifier

After training the CNN model for 30 epochs on the CIFAR-10 dataset, the model was evaluated using standard classification metrics on the test set.

Evaluation Metrics

- **Accuracy:** 74.2%
- **Loss (Cross Entropy):** 0.68
- **Precision (macro avg):** 73.1%
- **Recall (macro avg):** 72.5%
- **F1 Score (macro avg):** 72.6%




Figure 8: Confusion matrix for the CNN model on CIFAR-10 test data.

Class-wise Performance

- **Best Performing Classes:** Truck, Bird
- **Most Confused Classes:** Cat vs. Dog, Deer vs. Horse

Observations

- Model convergence occurred by around the 25th epoch.
- Data augmentation played a vital role in improving generalisation.

- Precision and recall were lower for semantically similar classes (e.g., cats vs dogs).
- Use of batch normalisation and dropout helped avoid overfitting.
- Training with a deeper architecture (e.g., ResNet) could further improve performance.

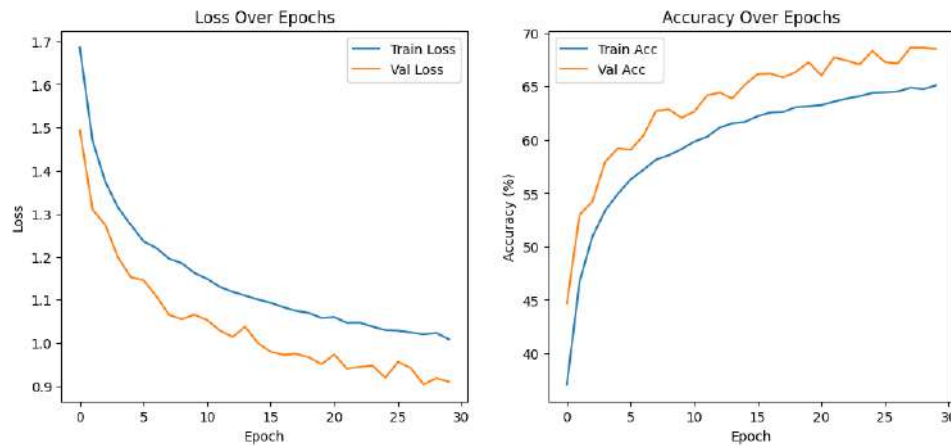


Figure 9: Training and validation accuracy and loss curves for the CNN model over 30 epochs. The steady decrease in loss and increase in accuracy indicate effective learning and minimal overfitting.

11 Conclusion

This internship represented a deep engagement with neural networks — both in theory and implementation — and allowed for an immersive experience across structured and unstructured data domains.

On the one hand, Artificial Neural Networks (ANNs) were trained and tested on a cardiac health dataset, demonstrating the feasibility of applying machine learning to healthcare tasks. Emphasis was placed not only on model performance but also on generalisation techniques such as dropout and batch normalisation, along with the choice of appropriate activation functions and loss metrics.

On the other hand, Convolutional Neural Networks (CNNs) were deployed for image classification tasks using the CIFAR-10 dataset. This allowed for an in-depth understanding of spatial hierarchies, convolutional layers, pooling operations, and the trade-offs involved in CNN depth, complexity, and performance. Regularisation via data augmentation, dropout, and batch norm proved crucial in achieving reasonable accuracy without overfitting.

While the results were not state-of-the-art, they were pedagogically significant — serving as a bridge between classroom instruction and applied machine learning. Challenges such as tuning hyperparameters, interpreting validation loss behaviour, and managing false positives/negatives in classification tasks brought clarity to the practical aspects of deep learning that are often abstracted in theoretical courses.

Perhaps most importantly, the internship reaffirmed the importance of experimentation, patience, and mathematical intuition in working with machine learning models. The neural network, for all its elegance, remains a computational approximation — and its behaviour, like all approximations, must be studied, scrutinised, and improved iteratively.

This report documents that effort — not as an endpoint, but as the beginning of more thoughtful and rigorous work in the field of deep learning.

Bibliography

References

- [1] CampusX, “100 Days of Deep Learning,” YouTube Playlist. https://www.youtube.com/playlist?list=PLKnIA16_Rmvbr7zKYQuBfsVkjoLcJgxHH
- [2] Krish Naik, “Deep Learning and Neural Networks,” YouTube Series. <https://youtu.be/d2kxUVwWwU?feature=shared>
- [3] NeuralNine, “Image Classification CNN in PyTorch,” *YouTube*, Available at: <https://youtu.be/CtzfbUwrYGI?feature=shared>
- [4] “PyTorch Documentation.” <https://pytorch.org/docs/stable/index.html>
- [5] “Google Colaboratory.” <https://colab.research.google.com/>
- [6] A. San Shuvo, “Cardiac Data (NHANES),” Kaggle Dataset. <https://www.kaggle.com/datasets/amithasanshuvo/cardiac-data-nhanes>
- [7] M. Elgendy, *Deep Learning for Vision Systems*, Manning Publications, 2020.
- [8] Unknown, “This figure shows three different activation functions graphs and its equations,” ResearchGate, 2017. Available at: https://www.researchgate.net/figure/This-figure-shows-three-different-activation-functions-graphs-and-its-equations-From_fig4_322537764
- [9] L. Poma, “Mastering Multi-Layer Perceptrons,” Medium, May 13, 2021. Available at: <https://medium.com/@lmpo/mastering-multi-layer-perceptrons-e7a82df6e844>
- [10] SaffronEdge, “Feedforward vs Backpropagation ANN,” LinkedIn, Feb 24, 2023. Available at: <https://www.linkedin.com/pulse/feedforward-vs-backpropagation-ann-saffronedge1/>